

Octave Compiler

Rob Vermaas

Department of Information & Computing Sciences
Utrecht University,
The Netherlands

May 4, 2005

Language

- ▶ GNU Octave is a high-level language, primarily intended for numerical computations
- ▶ Open-source clone of MATLAB
- ▶ Interpreter
- ▶ Big set of (internal) library functions

Why?

- ▶ Heavily used for prototyping and simulations
 - ▶ Developers like the interactiveness of the interpreter
 - ▶ Only prototyping, real implementation in low-level language
- ▶ High-level language with high-level data structures
 - ▶ Programs are getting bigger
 - ▶ No need to worry about declarations
 - ▶ Easy, convenient syntax

But, this comes at a cost.

Why can Octave code be so slow?

Consider the following code,

```
b = c = zeros(10);  
  
for x = 1:10  
    a(x) = b(x)*c(x) ;  
end
```

What happens at the assignment in the loop?

- ▶ Disambiguation
 - ▶ Distinguish function calls from array subscripts
- ▶ Type checking
 - ▶ Overloaded operators
- ▶ Bounds checking
- ▶ Reshaping

Octave compiler structure

- ▶ Octave Frontend
- ▶ Octave Optimizer
- ▶ Octave Type Inferencer
- ▶ Octave to C++ Backend

- ▶ Tools

- ▶ Parsing and packing Octave code
- ▶ Desugaring
- ▶ Format checker
- ▶ Pretty-printer

- ▶ Library

- ▶ Generic Forward Propagation Strategy
- ▶ Stratego-bindings for Octave interpreter-related functions
- ▶ An SDF syntax definition for Octave

Parsing Octave

- ▶ No proper description of language
 - ▶ Manual dates from 1997!
- ▶ Old situation
 - ▶ Adaptation of parser of interpreter which needed complete evaluation
- ▶ New situation
 - ▶ Also based on interpreter
 - ▶ Bison/Flex based
 - ▶ Wrapper that outputs an ATerm
 - ▶ Foreign function interface (primitives)
- ▶ Have we abandoned SDF ?
 - ▶ No! We need SDF! We need concrete syntax!

Octave Optimizer

- ▶ Source-to-source data-flow transformations
- ▶ Classic approach
 - ▶ Build control flow graph and put in single-static-assignment form
 - ▶ Perform optimization
 - ▶ Convert back to 'normal' source code
- ▶ Our approach
 - ▶ Stay close to source, work on abstract syntax tree
 - ▶ Use generic data-flow-specification to perform optimizations
- ▶ Advantages
 - ▶ All information is available
 - ▶ Analysis and transformation in one step
 - ▶ Easy combination of optimizations
- ▶ Disadvantages
 - ▶ Optimizations using SSA can probably easier be reused for other languages (the optimization step on SSA-part)
 - ▶ Use general/known algorithms

Currently, `octavec` can perform the following optimizations,

- ▶ Constant folding
- ▶ Partial Evaluation
 - ▶ Constant Propagation
 - ▶ Copy Propagation
 - ▶ Common Subexpression Elimination
 - ▶ Function specialization (value-based)
- ▶ Dead Code Elimination
- ▶ Vectorization

Not much changed here, except rewriting the specifications to new style dataflow strategies.

- ▶ Prototyping in Tiger
- ▶ Easy transition from Tiger to Octave
- ▶ GFPS, Constant/Copy propagation, CSE within a day
- ▶ Basically only encode scopes of the language

Constant propagation

- ▶ Discover values that are constant on all possible executions of a program and to propagate these values through the program
- ▶ Best combined with constant folding and branch elimination

Constant propagation

- ▶ Discover values that are constant on all possible executions of a program and to propagate these values through the program
- ▶ Best combined with constant folding and branch elimination

```
x = 1  
  
if x  
    y = x+3  
else  
    y = 2  
end  
  
disp(y)
```

Constant propagation

- ▶ Discover values that are constant on all possible executions of a program and to propagate these values through the program
- ▶ Best combined with constant folding and branch elimination

```
x = 1  
  
if x  
    y = x+3  
else  
    y = 2  
end  
  
disp(y)
```

```
x = 1  
y = 4  
  
disp(4)
```

Constant propagation, implementation

```
strategies
  octave-prop-const =
    forward-prop(
      prop-const-transform(octave-prop-const)
      , prop-const-before(octave-prop-const)
      , octave-prop-const
      , prop-const-after
      | ["PropConst"], [], []
    )

  prop-const-transform(recur) =
    ElimIf; recur
    <& [| while <recur> do <id> end ||]; ElimWhile; recur

  prop-const-before(recur) =
    [| <id>[<*recur>] = <id> ||

  prop-const-after =
    prop-const-assign
    <& PropConst <& EvalBinOp <& EvalRelOp

  prop-const-assign =
    ?[ x = e ]
    ; where( <is-value> e )
    ; rules( PropConst.x : [| x || -> [| e || depends on [(x,x)] ] )
```

- ▶ Combination of
 - ▶ Value-based function specialization
 - ▶ Constant folding
 - ▶ Constant propagation
 - ▶ Copy propagation
 - ▶ Common subexpression elimination
- ▶ For each call, determine static arguments and specialize the called function for these arguments, and apply the given optimizations

Partial evaluation, implementation

```
strategies
  partial-eval =
    forward-prop(
      prop-const-transform(partial-eval) + pe-transform
      , prop-const-before(partial-eval)
      , partial-eval
      , pe-after
      | ["PropConst","CopyProp","CSE"], [], []
    )

  pe-transform =
    ?[ f(a*) ]
    ; specialize-call(specialize-function-by-value,specialization-facts-by-value | <get-nargout>)

  pe-after =
    try(copy-prop-after); try(prop-const-after); try(cse-after)
```

Partial evaluation, implementation

```
strategies
  partial-eval =
    forward-prop(
      prop-const-transform(partial-eval) + pe-transform
      , prop-const-before(partial-eval)
      , partial-eval
      , pe-after
      | ["PropConst","CopyProp","CSE"], [], []
    )

  pe-transform =
    ?[ f(a*) ]
  ; specialize-call(specialize-function-by-value,specialization-facts-by-value | <get-nargout>)

  pe-after =
    try(copy-prop-after); try(prop-const-after); try(cse-after)
```

Okay, cheated a bit by leaving out the function specialization part.

```
$ wc -l /pkg/src/octave/octave-opt/eval/Octave-PE.str
78 /pkg/src/octave/octave-opt/eval/Octave-PE.str
```


Why loop vectorization?

Remember,

```
b = c = zeros(10);  
  
for x = 1:10  
    a(x) = b(x)*c(x) ;  
end
```

- ▶ Allows for parallelization
- ▶ Eliminates shape checks
- ▶ Eliminates reshapes
- ▶ Loop vectorization makes Octave code hard to read, therefore let compiler do it

Loop vectorization example

Original

```
function img = mono(r,g,b)
    [n,m] = size(r);
    img = zeros(n,m);

    for i = 1:n
        for j = 1:m
            img(i,j) = 0.3*r(i,j) + 0.6*g(i,j) + 0.1*b(i,j);
        end
    end
end
```

Loop vectorization example

Original

```
function img = mono(r,g,b)
    [n,m] = size(r);
    img = zeros(n,m);

    for i = 1:n
        for j = 1:m
            img(i,j) = 0.3*r(i,j) + 0.6*g(i,j) + 0.1*b(i,j);
        end
    end
end
```

Vectorized

```
function img = mono_vec(r,g,b)
    [n,m] = size(r);
    img = zeros(n,m);

    img(1:n,1:m) = 0.3*r(1:n,1:m) + 0.6*g(1:n,1:m) + 0.1*b(1:n,1:m);
end
```

- ▶ Type inferencing
 - ▶ Needed for efficient mapping in a backend
 - ▶ Possibilities for optimizations
 - ▶ Mainly a forward dataflow problem, therefore (again) GFPS
- ▶ Shape inference
 - ▶ Shape information is important for eliminating bound checks and reshaping
- ▶ Evaluating function calls and removing run-time checks

- ▶ Type inferencing
 - ▶ Needed for efficient mapping in a backend
 - ▶ Possibilities for optimizations
 - ▶ Mainly a forward dataflow problem, therefore (again) GFPS
- ▶ Shape inference
 - ▶ Shape information is important for eliminating bound checks and reshaping
- ▶ Evaluating function calls and removing run-time checks

```
x = [ 1, 2, 3, 4]  
[n, m] = size(x)
```

- ▶ Type inferencing
 - ▶ Needed for efficient mapping in a backend
 - ▶ Possibilities for optimizations
 - ▶ Mainly a forward dataflow problem, therefore (again) GFPS
- ▶ Shape inference
 - ▶ Shape information is important for eliminating bound checks and reshaping
- ▶ Evaluating function calls and removing run-time checks

```
x = [ 1, 2, 3, 4]
[n, m] = size(x)
```

```
x = [ 1, 2, 3, 4]
n = 1
m = 4
```

- ▶ Type inferencing
 - ▶ Needed for efficient mapping in a backend
 - ▶ Possibilities for optimizations
 - ▶ Mainly a forward dataflow problem, therefore (again) GFPS
- ▶ Shape inference
 - ▶ Shape information is important for eliminating bound checks and reshaping
- ▶ Evaluating function calls and removing run-time checks

```
x = [ 1, 2, 3, 4]
[n, m] = size(x)
```

```
x = [ 1, 2, 3, 4]
n = 1
m = 4
```

```
x = [1, 2, 3, 4]
if !ismatrix(x)
    printf("NO MATRIX!")
end
```

- ▶ Type inferencing
 - ▶ Needed for efficient mapping in a backend
 - ▶ Possibilities for optimizations
 - ▶ Mainly a forward dataflow problem, therefore (again) GFPS
- ▶ Shape inference
 - ▶ Shape information is important for eliminating bound checks and reshaping
- ▶ Evaluating function calls and removing run-time checks

```
x = [ 1, 2, 3, 4]
[n, m] = size(x)
```

```
x = [ 1, 2, 3, 4]
n = 1
m = 4
```

```
x = [1, 2, 3, 4]
if !ismatrix(x)
    printf("NO MATRIX!")
end
```

```
x = [1, 2, 3, 4]
```


User input

- ▶ Add type information for internal (non-user) octave functions
- ▶ Steer the type inference process
- ▶ Dynamic rules to interpret these rules

```
definitions
  constants
    pi      := float
    ...
  variables
    warn_fortran_indexing := int
    warn_num_to_str := int
    ...
  functions
    ...
    time      :: -> float
    fnmatch :: string, string -> int
    ...
```

- ▶ Embedded in Octave in Stratego

- ▶ Generates C++ code
 - ▶ Stand-alone application
 - ▶ Dynamically loaded function (.oct)
- ▶ Links against `liboctave`
- ▶ Fallback to dynamic typing (`liboctinterp`)

Example (Dynamically typed, DLD)

```
function x = factorial(n)
    x = 1;
    while(n>1)
        x = x * n;
        n = n - 1;
    end
end
```

Example (Dynamically typed, DLD)

```
function x = factorial(n)
    x = 1;
    while(n>1)
        x = x * n;
        n = n - 1;
    end
end
```

Dynamically loaded function (to generate .oct file)

```
DEFUN_DLD (factorial, args, nargout, "")
{
    octave_value_list c_0;
    octave_value x;
    octave_value n;
    n = args(0);
    x = 1;
    while ( do_binary_op(octave_value::op_gt, n, 1).all().all().bool_array_value()(0) )
    {
        x = do_binary_op(octave_value::op_mul, x, n);
        n = do_binary_op(octave_value::op_sub, n, 1);
    }
    c_0(0) = x;
    return(c_0);
}
```

Example (Whole program compilation, TI)

```
function x = factorial(n)
    x = 1;
    while(n>1)
        x = x * n;
        n = n - 1;
    end
end
```

```
factorial(40)
```

Example (Whole program compilation, TI)

```
function x = factorial(n)
    x = 1;
    while(n>1)
        x = x * n;
        n = n - 1;
    end
end
```

```
factorial(40)
```

```
void d_0__a_0 ()
{
    double f;
    f = factorial__n_int(40);
}
double factorial__n_int (double n)
{
    double x;
    x = 1;
    while ( (n > 1) )
    {
        x = (x * n);
        n = (n - 1);
    }
    return(x);
}
```

Future work

a lot...